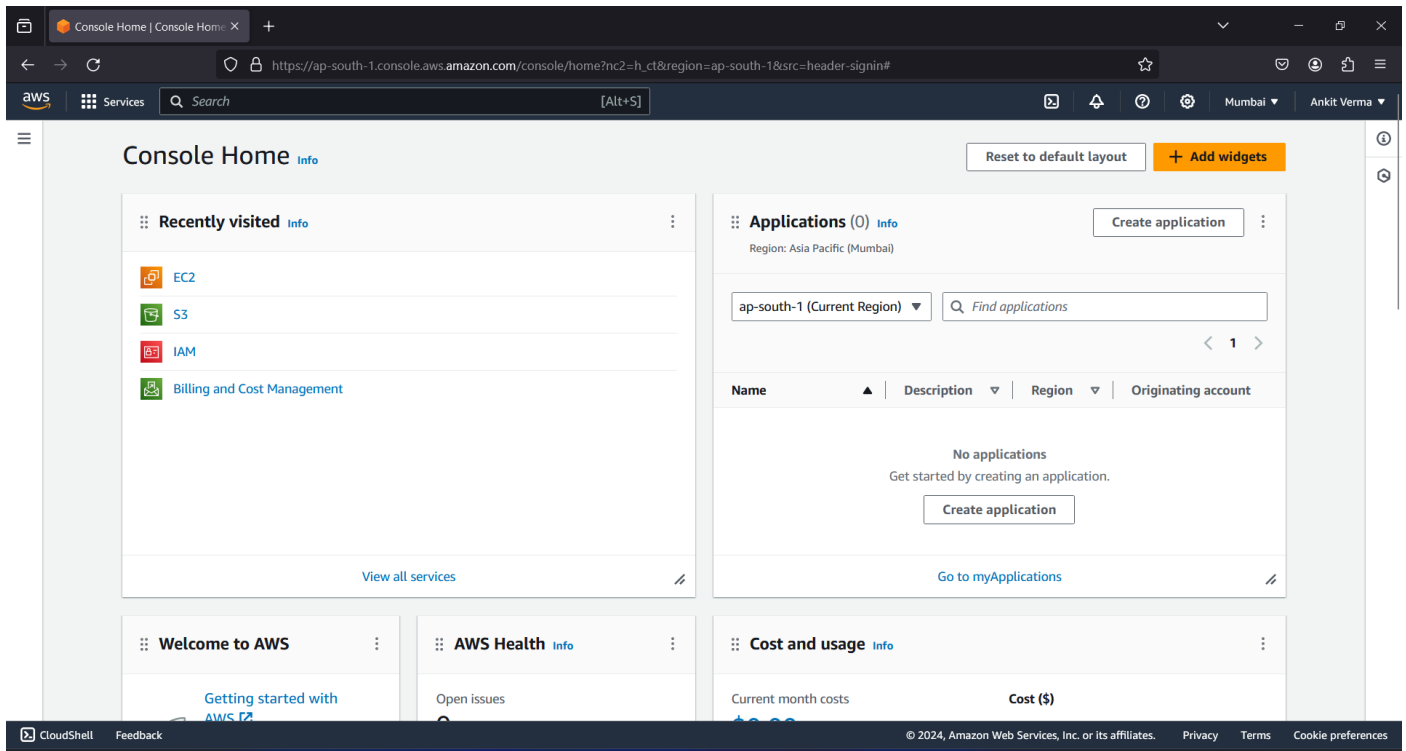


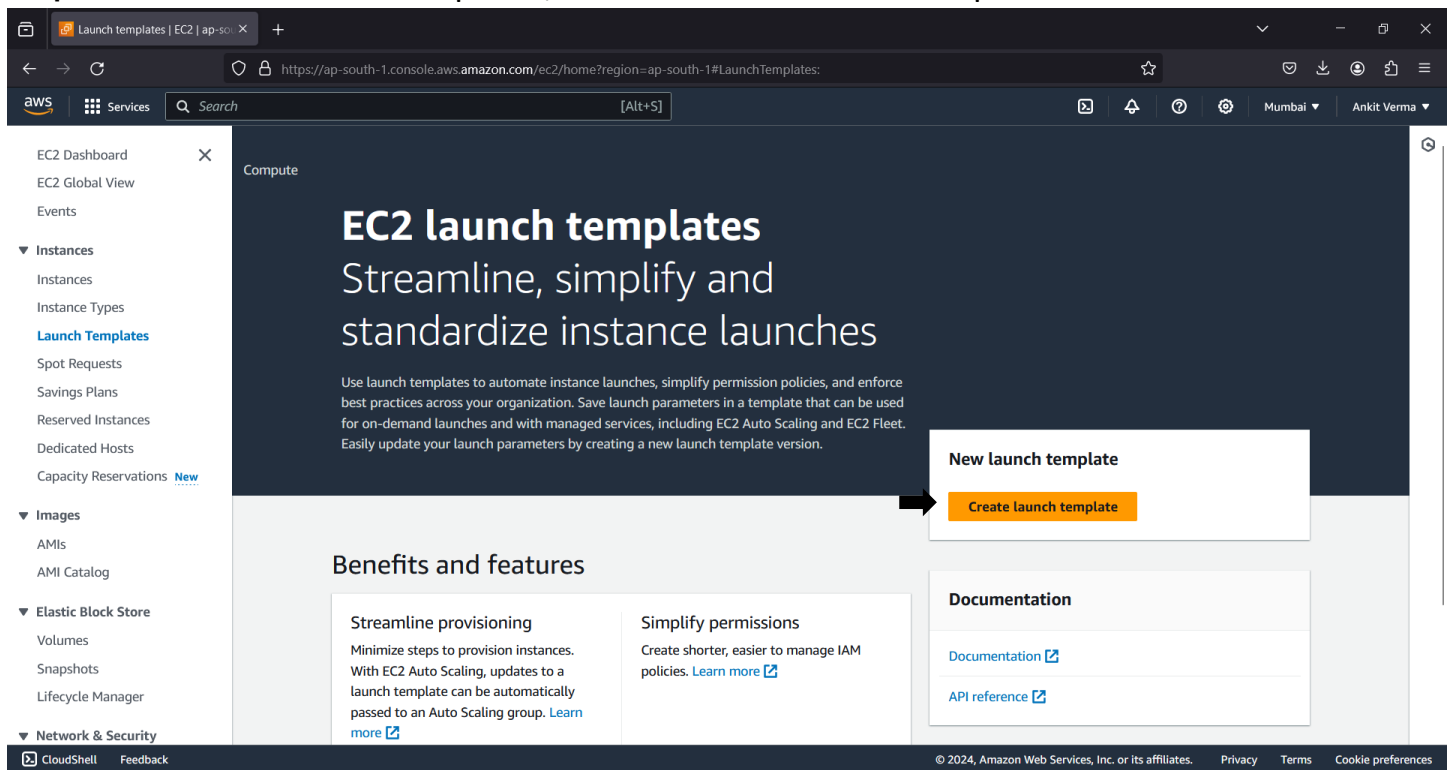
Assignment: 11

Problem Statement: Build scaling plans in AWS that balances the load on different EC2 instances.

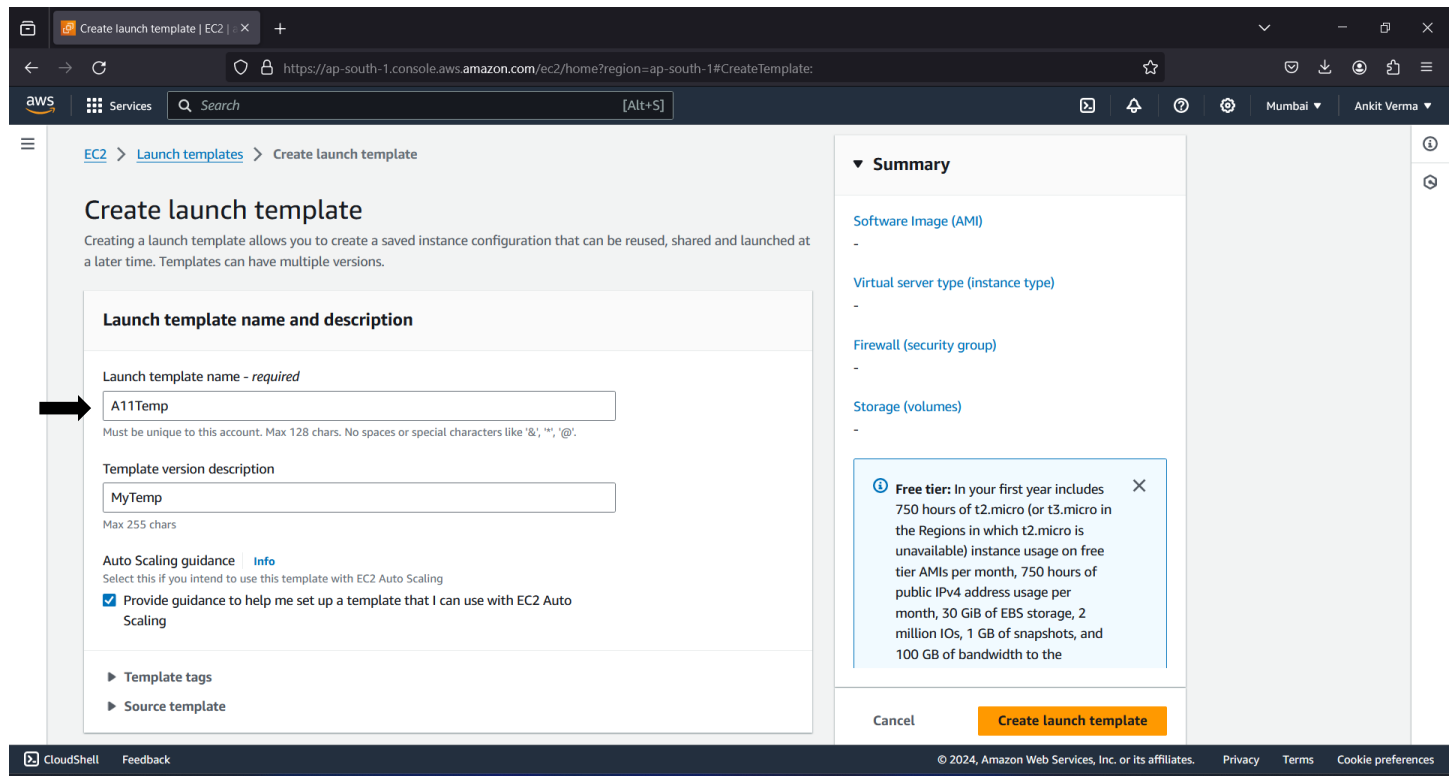
Step 1: Select EC2.



Step 2: Under the Launch Templates, click on Create Launch Template



Step 3: Give a name and description to the template. Check the Auto Scaling Guidance checkbox.



Create launch template | EC2 | +

https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateTemplate:

Services Search [Alt+S]

EC2 > Launch templates > Create launch template

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

A11Temp

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', "'", '@'.

Template version description

MyTemp

Max 255 chars

Auto Scaling guidance [Info](#)

Select this if you intend to use this template with EC2 Auto Scaling

☒ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

Template tags

Source template

Summary

Software Image (AMI)

-

Virtual server type (instance type)

-

Firewall (security group)

-

Storage (volumes)

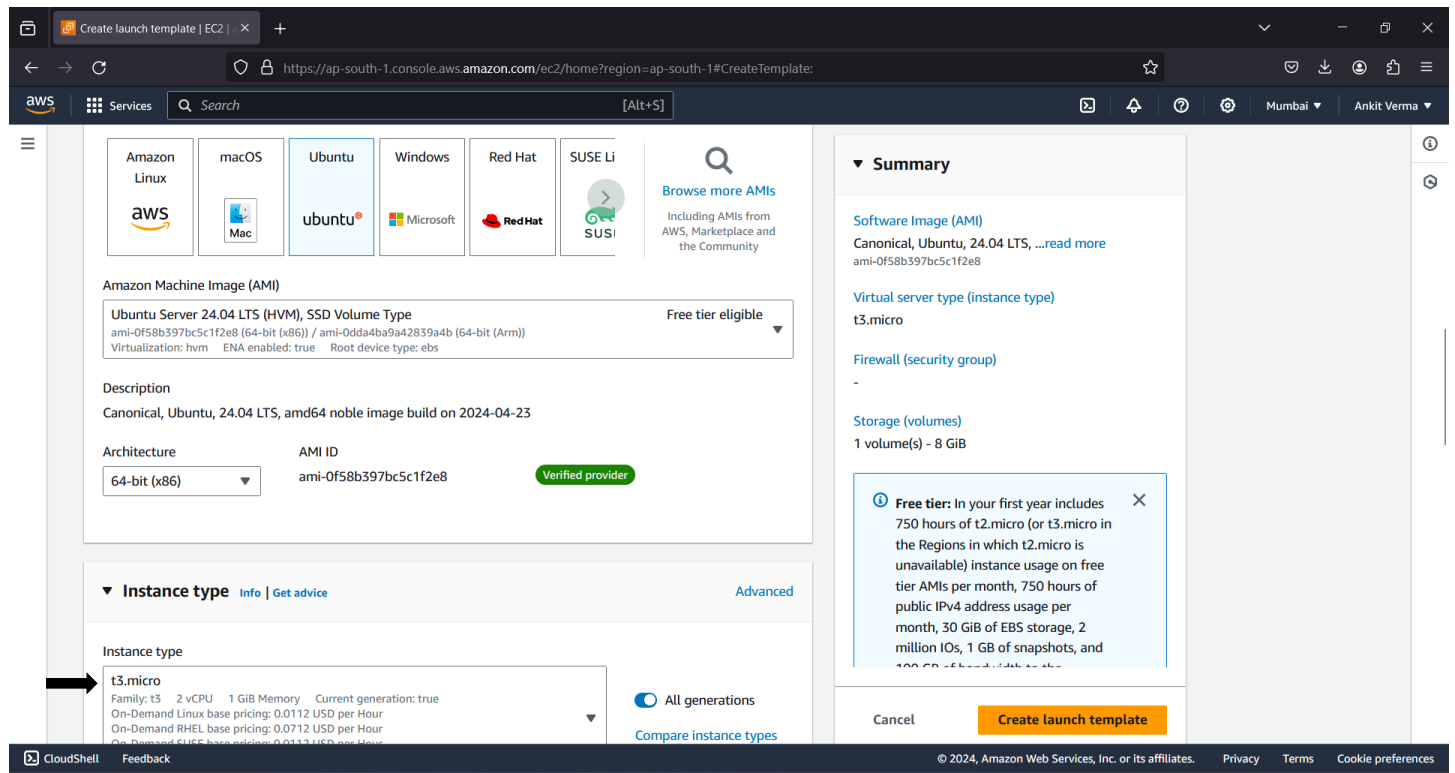
-

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the

Cancel Create launch template

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 4: Select Ubuntu & under instance type, select t3.micro .



Create launch template | EC2 | +

https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateTemplate:

Services Search [Alt+S]

Amazon Linux macOS **Ubuntu** Windows Red Hat SUSE Linux

aws Mac ubuntu Microsoft Red Hat SUSI

[Browse more AMIs](#)

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Ubuntu Server 24.04 LTS (HVM), SSD Volume Type

ami-0f58b397bc5c1f2e8 (64-bit (x86)) / ami-0dda4ba9a42839a4b (64-bit (Arm))

Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible

Description

Canonical, Ubuntu, 24.04 LTS, amd64 noble image build on 2024-04-23

Architecture

64-bit (x86)

AMI ID

ami-0f58b397bc5c1f2e8

Verified provider

Instance type

[Info](#) [Get advice](#) [Advanced](#)

Instance type

t3.micro

Family: t3 2 vCPU 1 GiB Memory Current generation: true

On-Demand Linux base pricing: 0.0112 USD per Hour

On-Demand RHEL base pricing: 0.0712 USD per Hour

On-Demand SUSE base pricing: 0.0112 USD per Hour

All generations

[Compare instance types](#)

Summary

Software Image (AMI)

Canonical, Ubuntu, 24.04 LTS, ...read more

ami-0f58b397bc5c1f2e8

Virtual server type (instance type)

t3.micro

Firewall (security group)

-

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the

Cancel Create launch template

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 5: Under key pair, select an existing key and select the user created Security group.

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name
key2 [Create new key pair](#)

Network settings [Info](#)

Subnet [Info](#)
Don't include in launch template [Create new subnet](#)

When you specify a subnet, a network interface is automatically added to your template.

Firewall (security groups) [Info](#)
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Select existing security group ☐ Create security group

Security groups [Info](#)
Select security groups
Assg10 sg-0e8da184778a43395 [Compare security group rules](#)

[Advanced network configuration](#)

Summary

Software Image (AMI)
Canonical, Ubuntu, 24.04 LTS, ...[read more](#)
ami-0f58b397bc5c1f2e8

Virtual server type (instance type)
t3.micro

Firewall (security group)
Assg10

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of S3 Standard storage.

[Cancel](#) [Create launch template](#)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 6: Expand the Advanced Details tab & Scroll down to the bottom, in the bash console type the following commands, Then click on Create Launch Template.

2

Allow tags in metadata [Info](#)
Don't include in launch template

User data - optional [Info](#)
Upload a file with your user data or enter it in the field.

[Choose file](#)

```
#!/bin/bash
apt-get update
apt-get install -y nginx
systemctl start nginx
systemctl enable nginx
apt-get install -y git
curl -sL https://deb.nodesource.com/setup_16.x | sudo -E bash -
apt-get install -y nodejs
git clone https://github.com/ankit11exe/Repo07.git
cd Repo07
npm install
node index.js
```

☐ User data has already been base64 encoded

Summary

Software Image (AMI)
Canonical, Ubuntu, 24.04 LTS, ...[read more](#)
ami-0f58b397bc5c1f2e8

Virtual server type (instance type)
t3.micro

Firewall (security group)
Assg10

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of S3 Standard storage.

[Cancel](#) [Create launch template](#)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 7: Click on Auto Scaling Group.

Success
Successfully created [A11Temp\(tt-01567bc1204b5c991\)](#).

Actions log

Next Steps

Launch an instance
With On-Demand Instances, you pay for compute capacity by the second (for Linux, with a minimum of 60 seconds) or by the hour (for all other operating systems) with no long-term commitments or upfront payments. Launch an On-Demand Instance from your launch template.
[Launch instance from this template](#)

Create an Auto Scaling group from your template
Amazon EC2 Auto Scaling helps you maintain application availability and allows you to scale your Amazon EC2 capacity up or down automatically according to conditions you define. You can use Auto Scaling to help ensure that you are running your desired number of Amazon EC2 instances during demand spikes to maintain performance and decrease capacity during lulls to reduce costs.
[Create Auto Scaling group](#)

Create Spot Fleet
A Spot Instance is an unused EC2 instance that is available for less than the On-Demand price. Because Spot Instances enable you to request unused EC2 instances at steep discounts, you can lower your Amazon EC2 costs significantly. The hourly price for a Spot Instance (of each instance type in each Availability Zone) is set by Amazon EC2, and adjusted gradually based on the long-term supply of and demand for Spot Instances. Spot instances are well-suited for data-analysis, batch jobs, background processing, and optional tasks.
[Create Spot Fleet](#)

Step 8: Click on Create Auto Scaling Group.

Amazon EC2 Auto Scaling helps maintain the availability of your applications

Auto Scaling groups are collections of Amazon EC2 instances that enable automatic scaling and fleet management features. These features help you maintain the health and availability of your applications.

Create Auto Scaling group

Get started with EC2 Auto Scaling by creating an Auto Scaling group.

[Create Auto Scaling group](#)

How it works

Auto Scaling group

Minimum size

Scale out as needed

Pricing

Amazon EC2 Auto Scaling features have no additional fees beyond the service fees for Amazon EC2, CloudWatch (for scaling policies), and the other AWS resources that you use. Visit the pricing page of each service to learn more.

Getting started

Step 9: Give a name and select the newly created Template. Then, click on Next.

The screenshot shows the 'Choose launch template' step in the AWS Management Console. The left sidebar lists steps 1 through 7. The main content area has a title 'Choose launch template' and a subtitle 'Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.' There are two main sections: 'Name' and 'Launch template'. The 'Name' section has a text input field with 'A11Scale' and a note that the name must be unique and no more than 255 characters. The 'Launch template' section has a dropdown menu with 'A11Temp' selected and a 'Create a launch template' link. A blue information box states: 'For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.' At the bottom, there are 'Cancel', 'Skip to review', 'Previous', and 'Next' buttons.

Step 1
Choose launch template

Step 2
Choose instance launch options

Step 3 - optional
Configure advanced options

Step 4 - optional
Configure group size and scaling

Step 5 - optional
Add notifications

Step 6 - optional
Add tags

Step 7
Review

Choose launch template Info

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.

Name

Auto Scaling group name
Enter a name to identify the group.

A11Scale

Must be unique to this account in the current Region and no more than 255 characters.

Launch template Info

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template
Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

A11Temp

[Create a launch template](#)

Version

Default (1)

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 10: In Network tab select all the available zones. Then click on next.

The screenshot shows the 'Network' step in the AWS Management Console. The left sidebar lists steps 6 and 7. The main content area has a title 'Network' and a subtitle 'For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.' There are two main sections: 'VPC' and 'Availability Zones and subnets'. The 'VPC' section has a dropdown menu with 'vpc-04d3f0c91f9f1eefc' selected and a 'Create a VPC' link. The 'Availability Zones and subnets' section has a dropdown menu with 'Select Availability Zones and subnets' selected and a 'Create a subnet' link. Below the dropdown, there are three subnets listed: 'ap-south-1a | subnet-01798752a060121b9', 'ap-south-1b | subnet-0e614382616965ac2', and 'ap-south-1c | subnet-079220daa9eeaaa02'. At the bottom, there are 'Cancel', 'Skip to review', 'Previous', and 'Next' buttons.

Step 6 - optional
[Add tags](#)

Step 7
[Review](#)

Network Info

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC
Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-04d3f0c91f9f1eefc

[Create a VPC](#)

Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

ap-south-1a | subnet-01798752a060121b9

ap-south-1b | subnet-0e614382616965ac2

ap-south-1c | subnet-079220daa9eeaaa02

[Create a subnet](#)

Cancel Skip to review Previous **Next**

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 11: Select Attach a new load balancer, select Application Load Balancer & select Internet Facing.

The screenshot shows the 'Create Auto Scaling group' page in the AWS Management Console, specifically Step 4: Configure advanced options. The page is titled 'Configure advanced options' and includes a sidebar with navigation links for Step 4 (optional), Step 5 (optional), Step 6 (optional), and Step 7. The main content area is divided into two sections. The first section, 'Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.', contains three radio button options: 'No load balancer', 'Attach to an existing load balancer', and 'Attach to a new load balancer'. The 'Attach to a new load balancer' option is selected. The second section, 'Attach to a new load balancer', contains three sub-sections: 'Load balancer type' with two radio button options, 'Application Load Balancer' (selected) and 'Network Load Balancer'; 'Load balancer name' with a text input field containing 'A11Scale-1'; and 'Load balancer scheme' with two radio button options, 'Internal' and 'Internet-facing' (selected). The footer of the page includes 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates.

Step 4 - optional
[Configure group size and scaling](#)

Step 5 - optional
[Add notifications](#)

Step 6 - optional
[Add tags](#)

Step 7
[Review](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☐ Attach to an existing load balancer
Choose from your existing load balancers.

☒ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to a new load balancer
Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type
Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, [visit the Load Balancing console](#).

☒ Application Load Balancer
HTTP, HTTPS

☐ Network Load Balancer
TCP, UDP, TLS

Load balancer name
Name cannot be changed after the load balancer is created.

A11Scale-1

Load balancer scheme
Scheme cannot be changed after the load balancer is created.

☐ Internal

☒ Internet-facing

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 12: Give the port no. 4000 & select Create a target group. Then select No VPC Lattice Service.

The screenshot shows the 'Create Auto Scaling group' page in the AWS Management Console, specifically Step 5: Listeners and routing. The page is titled 'Listeners and routing' and includes a sidebar with navigation links for Step 4 (optional), Step 5 (optional), Step 6 (optional), and Step 7. The main content area is divided into two sections. The first section, 'Listeners and routing', contains three sub-sections: 'Protocol' with a dropdown menu set to 'HTTP'; 'Port' with a text input field containing '4000'; and 'Default routing (forward to)' with a dropdown menu set to 'Create a target group'. The second section, 'New target group name', contains a text input field containing 'A11Scale-1'. The third section, 'Tags - optional', contains a button 'Add tag' and a note '50 remaining'. The fourth section, 'VPC Lattice integration options', contains a sub-section 'Select VPC Lattice service to attach' with two radio button options: 'No VPC Lattice service' (selected) and 'Attach to VPC Lattice service'. The footer of the page includes 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates.

Listeners and routing
If you require secure listeners, or multiple listeners, you can configure them from the [Load Balancing console](#) after your load balancer is created.

Protocol Port Default routing (forward to)

HTTP 4000 Create a target group

New target group name
An instance target group with default settings will be created.

A11Scale-1

Tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them.

[Add tag](#)

50 remaining

VPC Lattice integration options [Info](#)
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.

Select VPC Lattice service to attach

☒ No VPC Lattice service
VPC Lattice will not manage your Auto Scaling group's network access and connectivity with other services.

☐ Attach to VPC Lattice service
Incoming requests associated with specified VPC Lattice target groups will be routed to your Auto Scaling group.

[Create new VPC Lattice service](#)

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 13: Check the Turn on Elastic Load Balancing Health checks checkbox. Give the Health Check Grace Period of 240 seconds. Click on NEXT.

Create Auto Scaling group | EC2

https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateAutoScalingGroup: 90%

Additional health check types - optional [Info](#)

☒ Turn on Elastic Load Balancing health checks [Recommended](#)

Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.

EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing. To avoid unexpected terminations, first verify the settings of these health checks in the [Load Balancer console](#).

☐ Turn on VPC Lattice health checks

VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.

Health check grace period [Info](#)

This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.

240 seconds

Additional settings

Monitoring [Info](#)

☐ Enable group metrics collection within CloudWatch

Default instance warmup [Info](#)

The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.

☐ Enable default instance warmup

Cancel Skip to review Previous Next

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 14: Under Desired capacity, give a size of 2. Under Scaling, give min capacity 2 & max capacity 3.

Create Auto Scaling group | EC2

https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateAutoScalingGroup: 90%

Step 2

[Choose instance launch options](#)

Step 3 - optional

[Configure advanced options](#)

Step 4 - optional

Configure group size and scaling

Step 5 - optional

[Add notifications](#)

Step 6 - optional

[Add tags](#)

Step 7

[Review](#)

group.

Group size [Info](#)

Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Units (number of instances)

Desired capacity

Specify your group size.

2

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

2

Equal or less than desired capacity

Max desired capacity

3

Equal or greater than desired capacity

Automatic scaling - optional

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 15: Select Target tracking scaling policy. And give the instance warmup time of 240 seconds. Then click on 'Next'.

The screenshot shows the 'Automatic scaling - optional' step in the AWS console. The page title is 'Create Auto Scaling group | EC2'. The URL is 'https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateAutoScalingGroup:'. The page has a sidebar with 'Services' and a search bar. The main content area is titled 'Automatic scaling - optional' and includes an 'Info' icon. It explains that you can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group. There are two radio button options: 'No scaling policies' (selected) and 'Target tracking scaling policy'. The 'Target tracking scaling policy' option is highlighted with a blue border and includes a description: 'Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.' Below the options, there is a text input field for 'Scaling policy name' with the value 'Target Tracking Policy'. There is a 'Metric type' dropdown menu with 'Average CPU utilization' selected. There is a 'Target value' input field with the value '50'. There is an 'Instance warmup' input field with the value '240' and the unit 'seconds'. There is a checkbox for 'Disable scale in to create only a scale-out policy' which is unchecked. The footer of the page includes 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates.

capacity

capacity

Automatic scaling - optional [Info](#)

Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ Target tracking scaling policy
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Target Tracking Policy

Metric type [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

Target value

50

Instance warmup [Info](#)

240 seconds

☐ Disable scale in to create only a scale-out policy

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 16: Click next.

The screenshot shows the 'Add notifications - optional' step in the AWS console. The page title is 'Create Auto Scaling group | EC2'. The URL is 'https://ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#CreateAutoScalingGroup:'. The page has a sidebar with 'Services' and a search bar. The main content area is titled 'Add notifications - optional' and includes an 'Info' icon. It explains that you can send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group. There is a text input field for 'Add notification'. At the bottom of the page, there are four buttons: 'Cancel', 'Skip to review', 'Previous', and 'Next'. The 'Next' button is highlighted in orange. The footer of the page includes 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates.

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
[Choose launch template](#)

Step 2
[Choose instance launch options](#)

Step 3 - optional
[Configure advanced options](#)

Step 4 - optional
[Configure group size and scaling](#)

Step 5 - optional
Add notifications

Step 6 - optional
[Add tags](#)

Step 7
[Review](#)

Add notifications - optional [Info](#)

Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.

Add notification

Cancel Skip to review Previous **Next**

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 17: Click on next.

The screenshot shows the 'Add tags - optional' step in the AWS Management Console. The left sidebar lists steps 1 through 7, with 'Add tags' selected. The main content area has the title 'Add tags - optional' and a description: 'Add tags to help you search, filter, and track your Auto Scaling group across AWS. You can also choose to automatically add these tags to instances when they are launched.' Below this is a blue information box stating: 'You can optionally choose to add tags to instances (and their attached EBS volumes) by specifying tags in your launch template. We recommend caution, however, because the tag values for instances from your launch template will be overridden if there are any duplicate keys specified for the Auto Scaling group.' Underneath is a 'Tags (0)' section with an 'Add tag' button and '50 remaining' text. At the bottom right are 'Cancel', 'Previous', and 'Next' buttons.

Step 18: Review all the data of the group to be created and click on Create Auto Scaling Group.

The screenshot shows the 'Review' step in the AWS Management Console. The left sidebar lists steps 1 through 7, with 'Review' selected. The main content area displays a summary of the configuration: 'Instance scale-in protection' with a toggle for 'Enable instance protection from scale in' (currently off); 'Step 5: Add notifications' with an 'Edit' button and 'No notifications' listed; and 'Step 6: Add tags' with an 'Edit' button and 'No tags' listed. At the bottom right are 'Cancel', 'Previous', and 'Create Auto Scaling group' buttons.

Step 19: After creating the scaling group, go back to Instances from the left side menu.

Auto Scaling groups (1) [Info](#)

Search your Auto Scaling groups

Name	Launch template/configuration	Instances	Status	Desired capacity	Min	Max	Availabil...
A11Scale	A11Temp Version Default	0	Updating capacity	2	2	3	ap-south-1...

0 Auto Scaling groups selected

Step 20: Since the capacity was given as 2, two instances are created. Now open any one of the instances by clicking on its id.

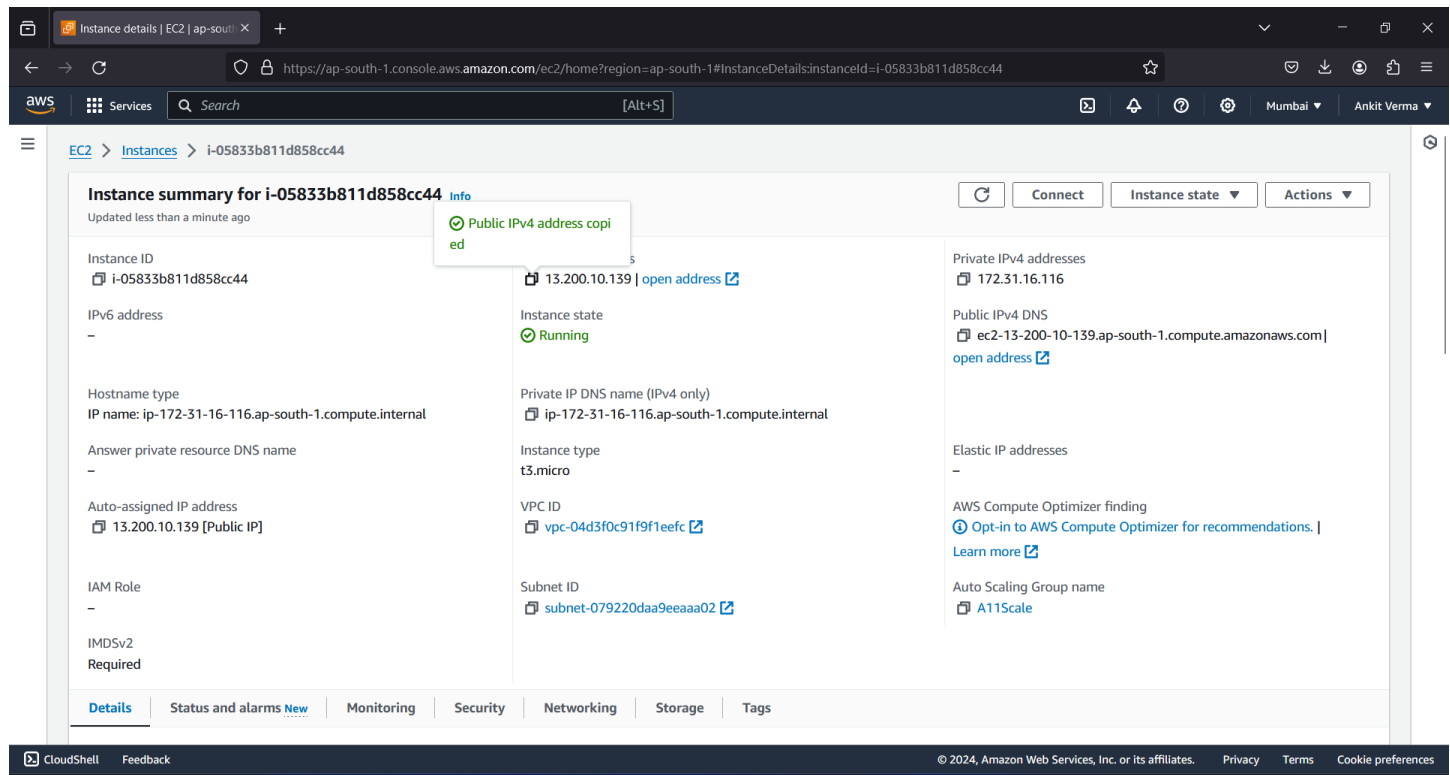
Instances (2) [Info](#)

Find Instance by attribute or tag (case-sensitive)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ..
	i-05833b811d858cc44	Running	t3.micro	2/2 checks passed	View alarms	ap-south-1c	ec2-13-200-10-139.ap...	13.200.10.13
	i-09a5667c74b0d2079	Running	t3.micro	2/2 checks passed	View alarms	ap-south-1b	ec2-3-110-56-114.ap-s...	3.110.56.114

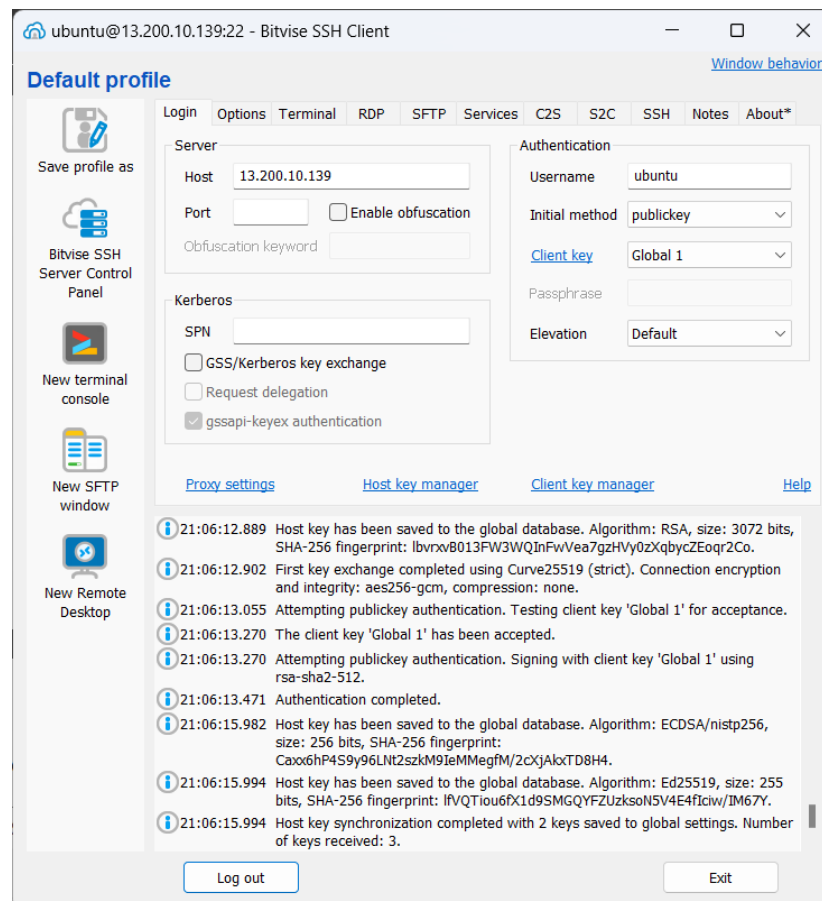
Select an instance

Step 21: Copy its Public IPv4 Address.



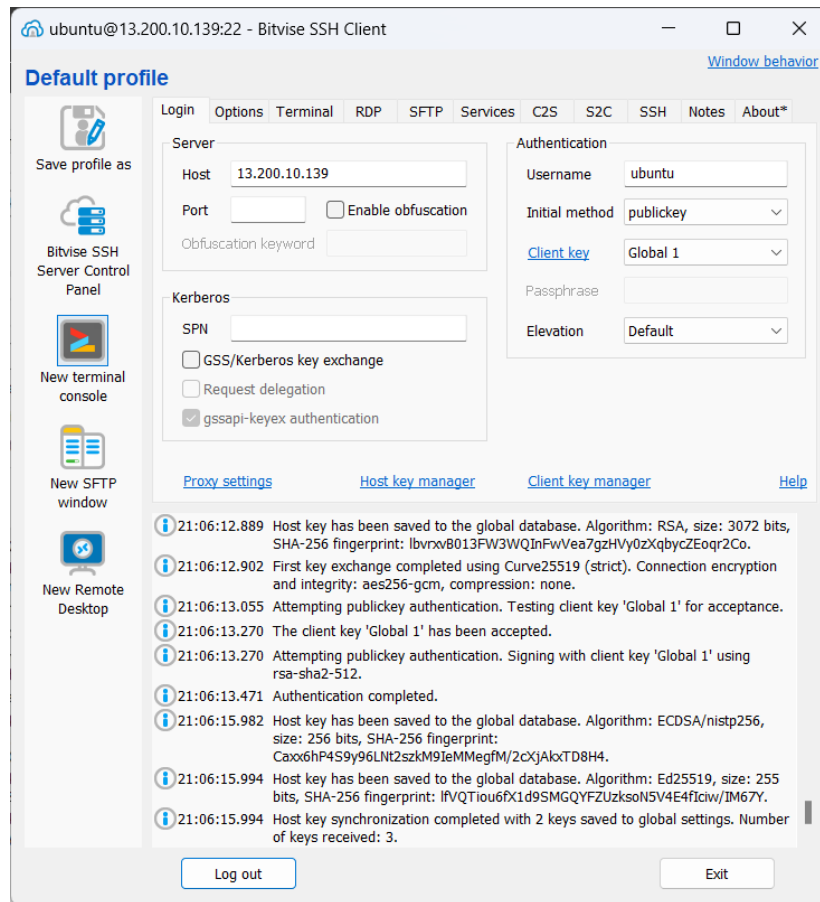
The screenshot shows the AWS Management Console for an EC2 instance. The instance ID is i-05833b811d858cc44. The instance is in the 'Running' state. The Public IPv4 address is 13.200.10.139, which is highlighted with a green circle and a tooltip that says 'Public IPv4 address copied'. The instance is running on a t3.micro instance type in the vpc-04d3f0c91f9f1eefc VPC, using the subnet-079220daa9eaaaa02 subnet. The Public IPv4 DNS is ec2-13-200-10-139.ap-south-1.compute.amazonaws.com.

Step 22: Paste the copied address and click on Log in.

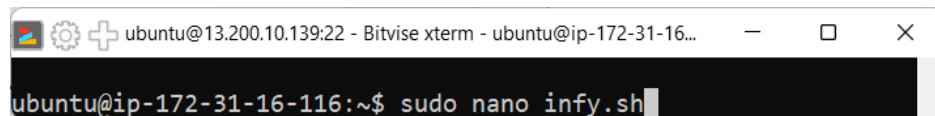


The screenshot shows the Bitvise SSH Client interface. The host is 13.200.10.139. The username is ubuntu. The initial method is publickey. The client key is Global 1. The passphrase is empty. The elevation is Default. The login process is shown in the log window, with messages indicating that the host key has been saved, the first key exchange is completed, and the client key 'Global 1' has been accepted.

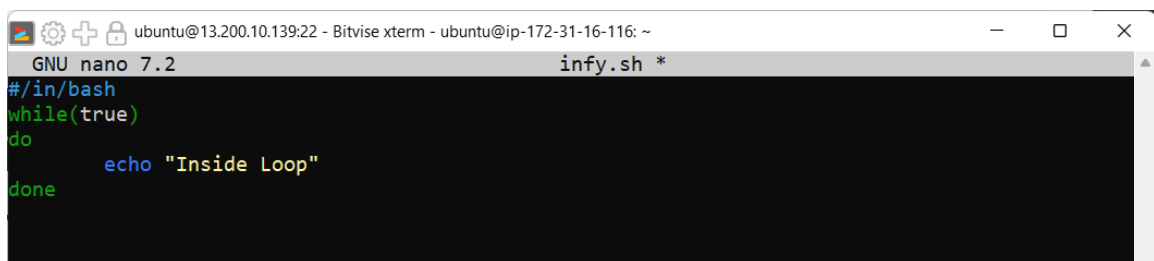
Step 23: Click on New Terminal Console.



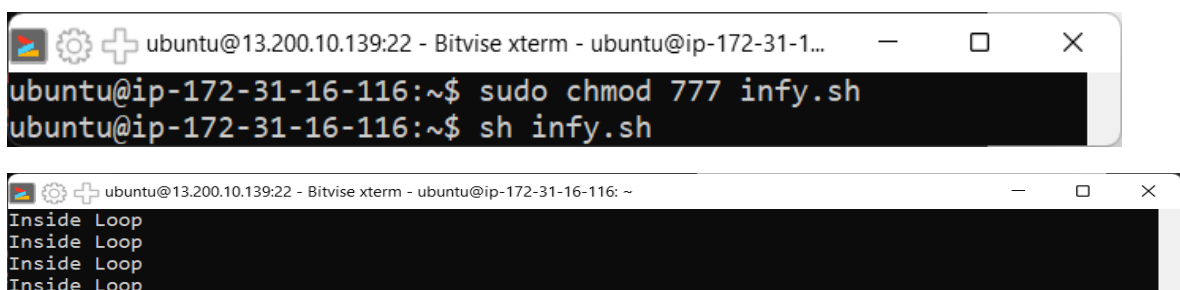
Step 24: Type the command:



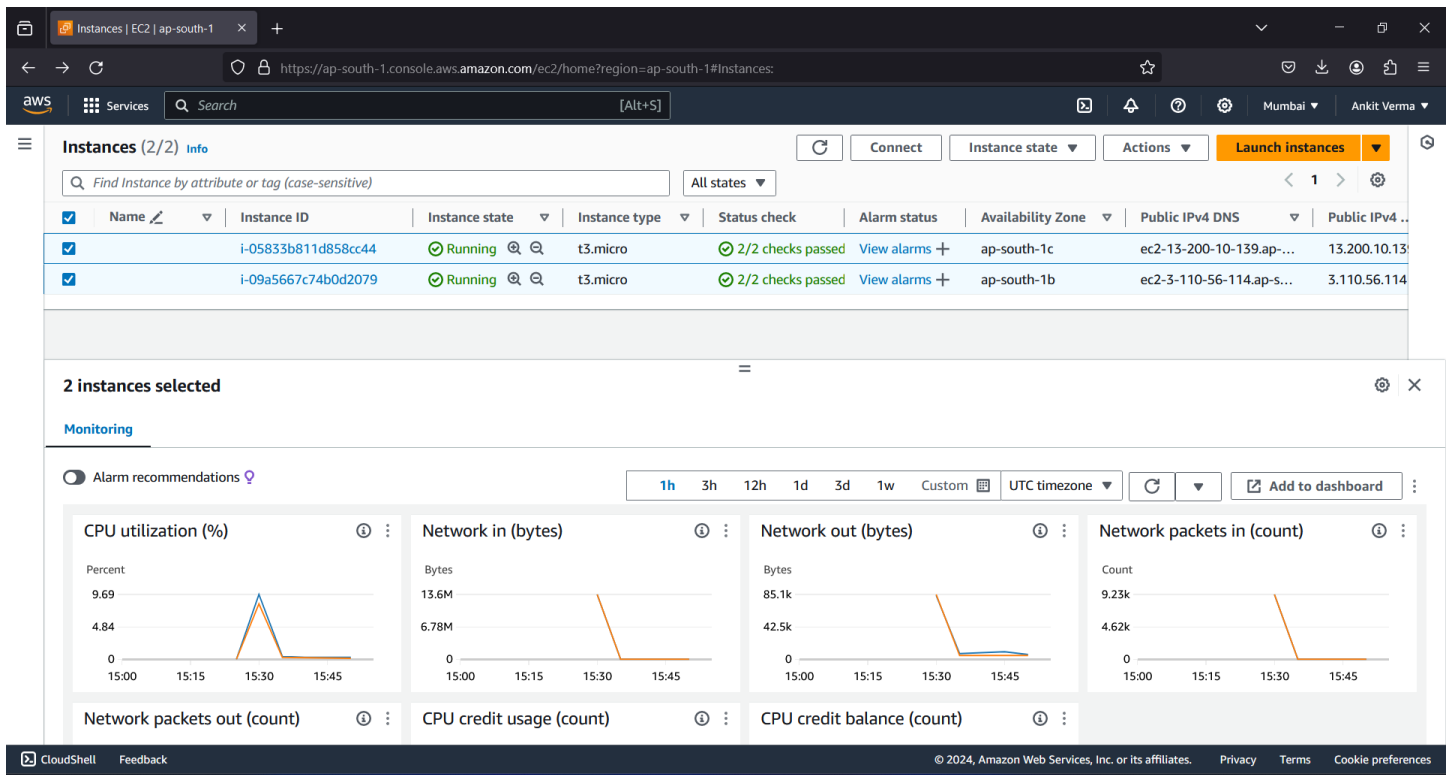
Step 25: Write the following code for an infinite loop in the infy.sh file.



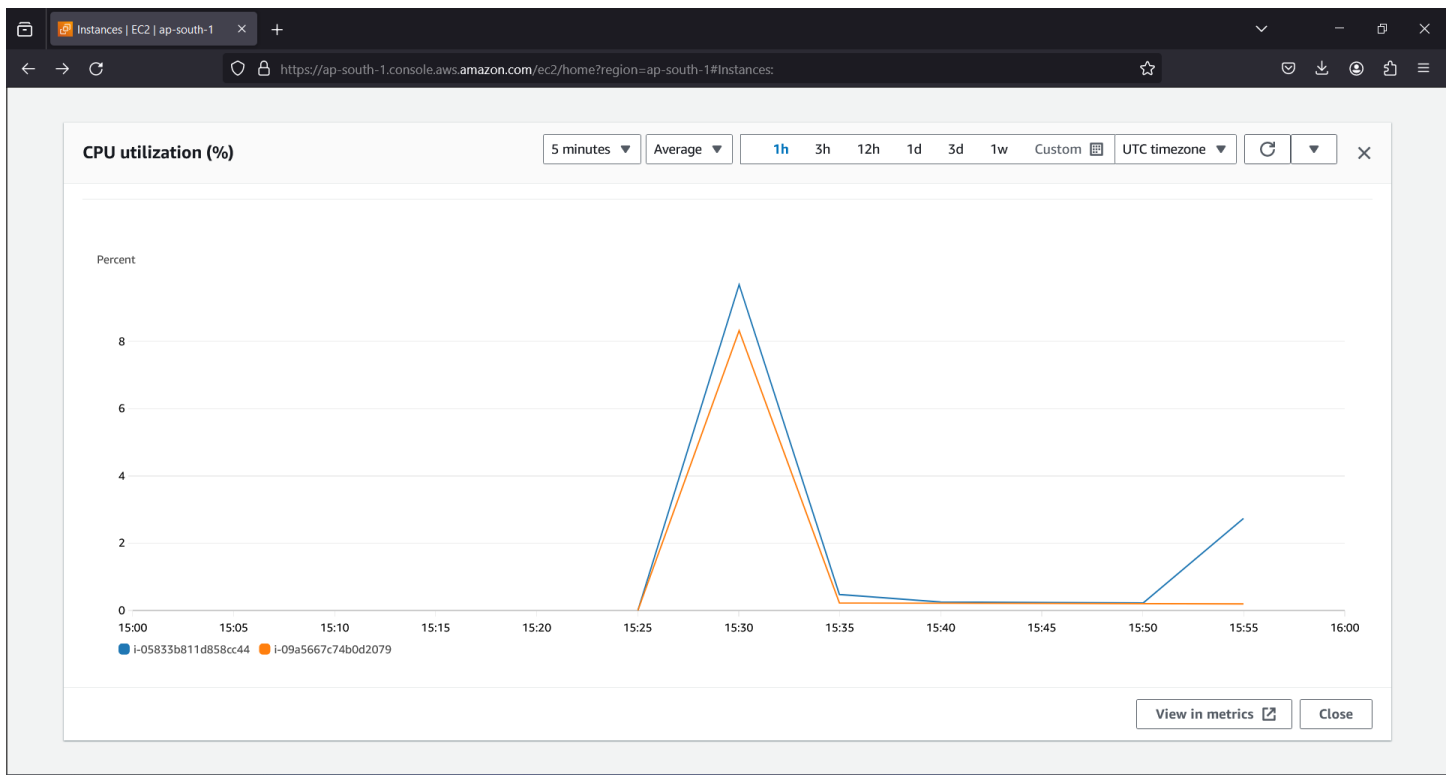
Step 26: Write the following commands in the terminal:



Step 27: Select both the instances, then under monitoring go to CPU utilization and enlarge it.



Step 28: The graph shows the CPU Utilization for both the instances.



When the CPU utilization exceed the limit for both the instances, a new instance will be created.